

Why Span Is Not Enough

Document Number: D4036R0
Date: 2026-02-28
Audience: LEWG
Reply-to: Vinnie Falco vinnie.falco@gmail.com

Table of Contents

Abstract

Revision History

R0: Feb 2026

Disclosure

Credit Where Due

1. `span<byte>`
2. `span<span<byte>>`
3. `range<span<byte>>`
4. `byte`

Boost.Asio

5. Six Ecosystems Already Arrived Here

6. The Final Straw

Almost There

7. Finally Correct

`void*` , Not `byte*`

A Buffer Sequence Is Distinct

What the Standard Needs

8. Side by Side

9. But

But this is standardizing Asio's types

But `vector<span<byte>>` is enough

But `mdspan` is enough

But `span<void>` is enough

But `span<byte>` is enough

Suggested Straw Poll

Acknowledgements

References

Abstract

C++ has bytes. A contiguous region of bytes needs a type. A sequence of such regions needs another. This paper examines the types that predictably come to mind, and their consequences.

Revision History

R0: Feb 2026

- Initial version.
-

Disclosure

The author maintains `Boost.Beast`^[7], a published HTTP and WebSocket library built on `Boost.Asio`^[1]'s buffer model, and develops `Capy`, `Corosio`, `Http`, `Beast2`, and `Burl`^[10] - libraries that define or consume buffer abstractions. The author published `P4003R0`^[8]. The author holds a neutral position on the Networking TS (changed from positive). This body of work creates a bias toward dedicated buffer types. Such types have costs: one more vocabulary type to learn, and interoperability friction with code that uses raw `span<byte>`.

Credit Where Due

`std::span` is a well-established vocabulary type. It turns a pointer and a size into a single thing. Perfectly. The vocabulary need is profound and this paper does not propose to diminish it.

The question is whether `span` is also the right vocabulary for I/O buffer descriptors.

1. `span<byte>`

Question. How do we represent a contiguous region of bytes?

Answer. `span<byte>`. A pointer and a size. That works.

However...

Platform I/O requires an array, not one region:

Platform	Descriptor	Used By
POSIX	<code>struct iovec</code>	<code>readv()</code> / <code>writew()</code>
POSIX	<code>struct msghdr</code>	<code>sendmsg()</code> / <code>recvmsg()</code>
Windows	<code>WSABUF</code>	<code>WSARecv()</code> / <code>WSASend()</code>
Windows	<code>FILE_SEGMENT_ELEMENT</code>	<code>ReadFileScatter()</code> / <code>WriteFileGather()</code>
Linux	<code>struct iovec</code>	<code>io_uring_prep_readv()</code> / <code>io_uring_prep_writew()</code>

`span<byte>` can describe one region. Wrap it in a one-element array and `readv()` accepts it. But I/O rarely involves a single contiguous region. A message has a header and a body. A protocol has framing and payload. Sending two regions with `write()` means two syscalls. Sending them with `writew()` means one - this is scatter/gather I/O. `span<byte>` is an insufficient type for representing an array of buffers.

2. `span<span<byte>>`

Question. How do we represent several such regions?

Answer. A span of spans. `span<span<byte>>`.

A single buffer is a view of someone's data. The bytes exist somewhere - in a `vector`, in a memory-mapped page, in a stack array. The buffer borrows them. Non-owning is natural. The data has a natural owner elsewhere.

A buffer sequence is different. Nobody "naturally" has an array of `span<byte>` objects lying around. The sequence is an assembled grouping - a data structure constructed to collect regions together. Making it non-owning means the grouping itself cannot be stored, returned, or passed across an asynchronous boundary.

```

class message {
    span<span<byte>> buffers_; // borrows... what?
};

span<span<byte>> prepare_message(span<byte> hdr, span<byte> body) {
    span<byte> bufs[] = { hdr, body };
    return { bufs }; // dangling
}

void start_send(socket& s, span<byte> hdr, span<byte> body) {
    span<byte> bufs[] = { hdr, body };
    s.async_send(span<span<byte>>(bufs), callback);
    // returns immediately; bufs destroyed; dangling
}

```

3. `range<span<byte>>`

Question. How do we own a collection of byte regions?

Answer. Use a range. `vector<span<byte>>`, `array<span<byte>, N>`, any range whose value type is `span<byte>`.

Ranges solve the ownership problem: a `vector` owns its elements.

Ranges create a byte consumption problem. Consider a JSON stream arriving in two chunks:

```

// chunk 1 [100 bytes]: {"name":"Alice","age":30}{"name":"B
// chunk 2 [100 bytes]: ob","age":25}...

range<span<byte>> input = { chunk1, chunk2 };

// parser finds first complete object: {"name":"Alice","age":30}
// that is 26 bytes - consume them
//
// views::drop(input, 1) drops all of chunk1 (100 bytes) - too much
// views::drop(input, 0) drops nothing - too little
// no standard range operation removes exactly 26 bytes

```

The parse boundary (26 bytes) does not align with the buffer boundary (100 bytes). Consuming 26 bytes means advancing chunk 1 by 26 bytes - 74 remain - without touching chunk 2. No range adaptor does this.

`std::ranges` ^[5] operates on elements. Parsing operates on bytes, not elements.

Incremental parsers with this need - JSON, XML, CSV, protobuf - go unserved.

4. `byte`

Question. What if we add byte-level algorithms to a range of `span<byte>` ?

Answer. The range is fine for ownership and iteration. The element type is not.

`span<byte>` already serves too many needs: serialization, cryptography, hashing, memory-mapped regions. If buffer sequences also use `span<byte>`, the type system cannot distinguish a buffer from any other byte span. A concept, an overload, or a constraint that separates “buffer in a sequence” from “hash input” or “encryption key” is impossible to write.

Boost.Asio

A separate type enables run-time safety checks:

Capability	Asio <code>mutable_buffer</code> ^[1]	<code>span<byte></code>
Implementation-defined members	Possible	Closed
Detect dangling after reallocation	Possible	No
Future diagnostic aids	Possible	No
Conditional debug callback	<code>BOOST_ASIO_ENABLE_BUFFER_DEBUGGING</code>	No

Each time `span<byte>` appears in a function signature, it loses the safety capability.

5. Six Ecosystems Already Arrived Here

Six I/O ecosystems, designed independently, all arrived at similar solutions:

Ecosystem	Buffer Type	Layout
POSIX	<code>iovec</code>	<code>void* + size_t</code>
Windows	<code>WSABUF</code>	<code>ULONG + char*</code>
Asio	<code>const_buffer</code> / <code>mutable_buffer</code>	<code>void const* + size_t</code> , with range concepts ^[1]
libuv	<code>uv_buf_t</code>	<code>char* + size_t</code> ^[2]

Ecosystem	Buffer Type	Layout
Go	<code>net.Buffers</code>	scatter/gather over <code>[]byte</code> ^[3]
.NET	<code>ReadOnlySequence<T></code>	linked list of discontinuous <code>Memory<T></code> segments ^[4]

Everybody converged on custom types independently.

6. The Final Straw

The committee already endorsed this principle.

P0298R3^[6] introduced `std::byte` because `unsigned char` performed triple duty. Neil MacIntosh wrote:

“these types perform a ‘triple duty’. Not only are they used for byte addressing, but also as arithmetic types, and as character types. This multiplicity of roles opens the door for programmer error”^[6]

“The key motivation here is to make byte a distinct type - to improve program safety by leveraging the type system.”^[6]

`unsigned char` had the right size and alignment. The committee added `std::byte` anyway - same size, same alignment, but no arithmetic, no implicit conversions. The generic type’s operations did not match the domain. The committee restricted the interface.

`span<byte>` performs double duty - general-purpose byte view and I/O buffer descriptor. A bespoke type restricts the interface to `data()` and `size()`. Same principle, one level of abstraction higher.

The precise fit is bespoke.

Almost There

`std::byte` kept the shift operators despite the stated goal of removing arithmetic. The principle was right. The execution left a gap.

7. Finally Correct

New buffer types give us the principled option. Only what we need: `data()` and `size()`.

`void*` , Not `byte*`

`void*` is maximally accepting and minimally permissive. Any pointer converts to it implicitly. The user must perform an explicit cast to go back. The asymmetry is by design.

Risk	<code>void*</code>	<code>byte*</code>	Cost
Requires <code>reinterpret_cast</code>	No	Yes	Invites superfluous casts
Dereferenceable	No	Yes	Invites accidental access
Pointer arithmetic	No	Yes	Invites accidental arithmetic
Assignable to <code>span<byte></code>	No	Yes	Invites full span API misuse
Promises byte-level meaning	No	Yes	Invites false type assertions
Contradicts type erasure	No	Yes	Invites type erasure violations
C++17 only	No	Yes	Disinvites C users

A Buffer Sequence Is Distinct

Buffer sequences are not served by existing concepts. They are a new concept.

What the Standard Needs

- A read-only byte region type (`void const*` + `size_t`)
- A writable byte region type (`void*` + `size_t`)
- Concepts for sequences of read-only and writable byte regions
- Algorithms: total byte count, byte-granular slicing, copy between buffer sequences

The types already exist:

```

class mutable_buffer {
    unsigned char* p_ = nullptr;
    std::size_t n_ = 0;
public:
    mutable_buffer() = default;
    mutable_buffer(mutable_buffer const&) = default;
    mutable_buffer& operator=(mutable_buffer const&) = default;
    constexpr mutable_buffer(void* data, std::size_t size) noexcept
        : p_(static_cast<unsigned char*>(data)), n_(size) { }
    constexpr void* data() const noexcept { return p_; }
    constexpr std::size_t size() const noexcept { return n_; }
};

class const_buffer {
    unsigned char const* p_ = nullptr;
    std::size_t n_ = 0;
public:
    const_buffer() = default;
    const_buffer(const_buffer const&) = default;
    const_buffer& operator=(const_buffer const& other) = default;
    constexpr const_buffer(void const* data, std::size_t size) noexcept
        : p_(static_cast<unsigned char const*>(data)), n_(size) { }
    constexpr const_buffer(mutable_buffer const& b) noexcept
        : p_(static_cast<unsigned char const*>(b.data())), n_(b.size()) { }
    constexpr void const* data() const noexcept { return p_; }
    constexpr std::size_t size() const noexcept { return n_; }
};

```

These are the **Networking TS^[9]** types.

8. Side by Side

Task	<code>span<byte></code>	<code>mutable_buffer</code>
Construct from vector	<code>span<byte>{reinterpret_cast<byte*>(v.data()), ...}</code>	<code>mutable_buffer{v.data(), v.size()}</code>
Consume N bytes	<code>buf = span<byte>{buf.data() + n, buf.size() - n}</code>	<code>buf += n</code>
Detect dangling	Requires ABI Break	<i>see-below</i>

Safety feature:

```
class mutable_buffer {  
    void* p_ = nullptr;  
    size_t n_ = 0;  
    void(*check_)() = nullptr;  
public:  
    void* data() const { if(check_) check_(); return p_; }  
    size_t size() const noexcept { return n_; }  
};
```

Smaller to write, safer to use, open to diagnostics.

9. But

But this is standardizing Asio's types

Yes. They earn their keep.

But `vector<span<byte>>` is enough

Users opt out of types which do not let them opt out of allocations.

But `mdspan` is enough

Buffer sequences only need one dimension. `mdspan`^[5] provides several.

But `span<void>` is enough

Even if `span<void>` were possible, what remains after removing the impossible is `data()` and `size()`. That is just a less-capable `mutable_buffer`.

But `span<byte>` is enough

`span<byte>` is also a less-capable `mutable_buffer`. It is `span<void>` with added harm.

Suggested Straw Poll

LEWG agrees that a contiguous byte region descriptor for I/O should be a dedicated type, not `span<byte>` .

Acknowledgements

The buffer model described here draws on twenty years of Asio's buffer sequence abstractions, due to Chris Kohlhoff.

References

1. **Boost.Asio** - Buffer types and buffer sequence requirements (Chris Kohlhoff).
https://www.boost.org/doc/libs/release/doc/html/boost_asio.html
2. **libuv** - `uv_buf_t` buffer type. <https://docs.libuv.org/en/v1.x/>
3. **Go standard library** - `net.Buffers` (<https://pkg.go.dev/net#Buffers>). <https://pkg.go.dev/>
4. **.NET System.IO.Pipelines** - `ReadOnlySequence<T>` . <https://learn.microsoft.com/en-us/dotnet/api/system.io.pipelines>
5. **C++ Working Draft** - `span` ([span.overview](#)), `mdspan` ([mdspan.overview](#)), ranges ([ranges](#)).
<https://eel.is/c++draft/>
6. **P0298R3** - A byte type definition (Neil MacIntosh). <https://wg21.link/p0298r3>
7. **Boost.Beast** - HTTP and WebSocket built on Boost.Asio (Vinnie Falco). <https://github.com/boostorg/beast>
8. **P4003R0** - (Vinnie Falco). <https://wg21.link/p4003r0>
9. **N4771** - Working Draft, C++ Extensions for Networking. <https://wg21.link/n4771>
10. **C++ Alliance** - Capy, Corosio, Http, Beast2, Burl (Vinnie Falco). <https://github.com/cppalliance>